

# RANDOMIZATION METHODS: BAGGING

-STATISTICAL MACHINE LEARNING-

Lecturer: Darren Homrighausen, PhD

# NOTATION

**REMINDER:** For either **classification** or **regression**, we produce **predictions** for a given feature vector  $X$

That is, we form

$$\hat{Y} = \hat{f}(X)$$

where

- $\hat{f}$  is some procedure formed with the **training data**  
(**EXAMPLE:**  $\hat{\beta}$  formed by least squares)
- The prediction  $\hat{Y}$  formed at a desired feature vector  $X$   
(**EXAMPLE:**  $\hat{Y} = X^\top \hat{\beta}$  formed by least squares)

# BAGGING

Many methods (trees included) tend to be designed to have lower bias but high variance

**HEURISTICALLY:** If we split the training data into two parts at random and fit a decision tree to each part, the results could be quite **different**

A low variance estimator would yield **similar** results if applied repeatedly to distinct data sets

(consider  $\hat{f}(X) = 0$  for all  $X$ )

**Bagging**, also known as **Bootstrap AGgregation**, is a general purpose procedure for reducing variance.

We'll use it specifically in the context of trees, but it can be applied more broadly.

# BAGGING: THE MAIN IDEA

Suppose we have  $n$  uncorrelated observations  $Z_1, \dots, Z_n$ , each with variance  $\sigma^2$ .

What is the variance of

$$\bar{Z} = \frac{1}{n} \sum_{i=1}^n Z_i?$$

# BAGGING: THE MAIN IDEA

Suppose we have  $n$  uncorrelated observations  $Z_1, \dots, Z_n$ , each with variance  $\sigma^2$ .

What is the variance of

$$\bar{Z} = \frac{1}{n} \sum_{i=1}^n Z_i?$$

ANSWER:  $\sigma^2/n$ .

More generally, if we have  $B$  separate (uncorrelated) training sets, we could form  $B$  separate model fits,

$$\hat{f}^1(X), \dots, \hat{f}^B(X)$$

Then average them:

$$\hat{f}_B(X) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(X)$$

# BAGGING: THE BOOTSTRAP PART

Of course, this isn't practical as having access to many training sets is unlikely.

We therefore turn to the **bootstrap** to simulate having many training sets.

The bootstrap is a widely applicable statistical tool that can be used to quantify uncertainty without Gaussian approximations.

Let's look at an example.

# Bootstrap detour

# BOOTSTRAP DETOUR

Suppose we are looking to invest in two financial instruments,  $X$  and  $Y$ . The return on these investments is random, but we still want to allocate our money in a risk minimizing way.

That is, for some  $\alpha \in (0, 1)$ , we want to minimize

$$\text{Var}(\alpha X + (1 - \alpha)Y)$$

The minimizing  $\alpha$  is:

$$\alpha_* = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}}$$

(Here,  $\sigma_{XY}$  is the **covariance** between  $X$  and  $Y$ )



# BOOTSTRAP DETOUR

We can estimate  $\alpha_*$  via a **plug-in** estimator

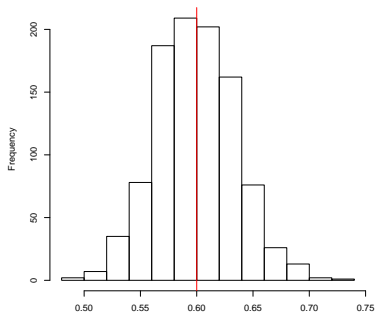
$$\hat{\alpha} = \frac{\hat{\sigma}_Y^2 - \hat{\sigma}_{XY}^2}{\hat{\sigma}_X^2 + \hat{\sigma}_Y^2 - 2\hat{\sigma}_{XY}^2}$$

Now that we have an estimator of  $\alpha_*$ , it would be nice to have an estimator of its **variability**.

In this case, computing a standard error is difficult.


# BOOTSTRAP DETOUR

Suppose for a moment that we can simulate a large number of draws (say 1000) of the data, which has actual value  $\alpha = 0.6$ . Then we could get estimates  $\hat{\alpha}_1, \dots, \hat{\alpha}_{1000}$ :



This is the **sampling distribution** of  $\hat{\alpha}$

# BOOTSTRAP DETOUR



The mean of all of these is:

$$\bar{\alpha} = \frac{1}{1000} \sum_{r=1}^{1000} \hat{\alpha}_r = 0.599,$$

which is very close to 0.6 (red line), and the standard error is

$$\sqrt{\frac{1}{1000 - 1} \sum_{r=1}^{1000} (\hat{\alpha}_r - \bar{\alpha})^2} = 0.079$$

## BOOTSTRAP DETOUR

The standard error of 0.035 gives a very good idea of the accuracy of  $\hat{\alpha}$  for a single sample.

Roughly speaking, for a new random sample, we expect

$$\hat{\alpha} \in (\alpha - 2 * 0.079, \alpha + 2 * 0.079) = (0.442, 0.758)$$

In practice, of course, we cannot use this procedure as it relies on being able to draw a large number of (independent) samples from the same distribution as our data.

This is where the **bootstrap** comes in.

We instead draw a large number of samples directly from our observed data. This sampling is done **with replacement**, which means that the same data point can be drawn multiple times.

# BOOTSTRAP DETOUR: SMALL EXAMPLE

Suppose we have data  $\mathcal{D} = (4.3, 3, 7.2, 6.9, 5.5)$ .

Then we can draw bootstrap samples, which might look like:

$$\mathcal{D}_1^* = (7.2, 4.3, 7.2, 5.5, 6.9)$$

$$\mathcal{D}_2^* = (6.9, 4.3, 3.0, 4.3, 6.9)$$

$$\vdots$$

$$\mathcal{D}_B^* = (4.3, 3.0, 3.0, 5.5, 6.9)$$

It turns out each of these  $\mathcal{D}_b^*$  have **very** similar properties as  $\mathcal{D}$

# BOOTSTRAP DETOUR: SMALL EXAMPLE

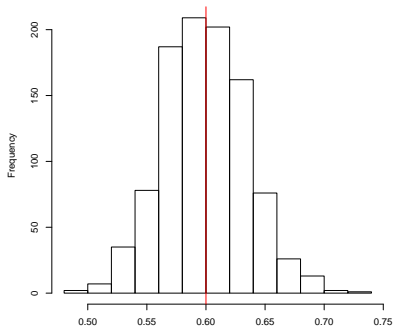
Now, we form the bootstrap mean:

$$\text{mean}_B = \frac{1}{B} \sum_{b=1}^B \hat{\alpha}_b^*$$

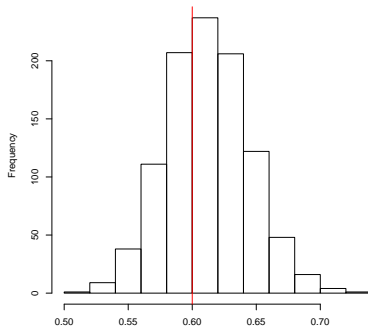
The bootstrap estimator of the standard error is:

$$\text{SE}_B = \sqrt{\frac{1}{B} \sum_{b=1}^B (\hat{\alpha}_b^* - \text{mean}_B)^2}$$

# BOOTSTRAP DETOUR



Sampling distribution of  $\hat{\alpha}$   
(impossible to form)



Bootstrap distribution of  $\hat{\alpha}$   
(possible to form)

# BOOTSTRAP: END DETOUR

## SUMMARY:

Suppose we want to get an idea of the sampling distribution of some statistic  $\hat{f}$  trained on  $\mathcal{D}$ .

Then we do the following: Fix a large number  $B$

( $B$  could be, say, 1000)

Then for each  $b = 1, \dots, B$

1. Form a new bootstrap draw from  $\mathcal{D}$ , call it  $\mathcal{D}^*$
2. Compute  $\hat{f}_b^*$  from  $\mathcal{D}^*$

Now, we can estimate the distribution of  $\hat{f}$  trained on  $\mathcal{D}$  by looking at the **distribution** of the  $B$  draws,  $\hat{f}_b^*$



# End detour

# BAGGING: THE BOOTSTRAP PART

Now, instead of having  $B$  separate training sets, we train on  $B$  **bootstrap** draws:

$$\hat{f}_1^*(X), \dots, \hat{f}_B^*(X)$$

and then average (i.e. **aggregate**) them:

$$\hat{f}_{\text{bag}}(X) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b^*(X)$$

This process is known as **Bagging**

# Bagging trees



# BAGGING TREES

The procedure for trees is the following

1. Choose a large number  $B$ .
2. For each  $b = 1, \dots, B$ , grow an unpruned tree on the  $b^{\text{th}}$  bootstrap draw from the data.
3. Average all these trees together.

Each tree, since it is unpruned, will have (low/high) variance and (low/high) bias

# BAGGING TREES

The procedure for trees is the following

1. Choose a large number  $B$ .
2. For each  $b = 1, \dots, B$ , grow an unpruned tree on the  $b^{\text{th}}$  bootstrap draw from the data.
3. Average all these trees together.

Each tree, since it is unpruned, will have (low/high) variance and (low/high) bias

Therefore averaging many trees results in an estimator that has lower variance and still low bias.

# BAGGING TREES IN CLASSIFICATION

For classification, there are a few sensible methods for aggregation

For each test observation  $X$ ,

- record the length  $B$  vector  $[\hat{f}_1^*(X), \dots, \hat{f}_B^*(X)]^\top$  and classify  $X$  via majority vote
- Average the length  $G$  probability vectors from each tree and choose the argmax

**WARNING:** One thing you definitely do not want to do is estimate probabilities via taking proportions of times  $X$  was classified to each class across the  $B$  trees

# Additional tree bagging topics

# BAGGING TREES

Now that we are growing a large number ( $B$ ) of random trees, we can't directly look at the **dendrogram**

This means we have sacrificed some interpretability for better performance

However, we do get some helpful information instead

- Mean decrease variable importance
- Out-of-Bag error estimation (OOB)
- Permutation variable importance
- Proximity

(Note that these ideas apply to bagging any low bias procedure, not just to unpruned trees)



# MEAN DECREASE VARIABLE IMPORTANCE

**Observation:** Every time a split of a node is made on a **feature**, the loss function is not increased

Hence, adding up the **loss decreases** for each feature over all trees gives an indication of feature importance

Intuitively an important feature is one that if split upon, it leads to a large reduction in the loss


# MEAN DECREASE VARIABLE IMPORTANCE

To recover some information, we can do the following:

1. For each of the  $B$  trees and each of the  $p$  features, we record the amount that the Gini index (or cross-entropy) is reduced by splitting on that feature
2. Report the average reduction over all  $B$  trees

This gives us an indication of the **importance** of a feature

# MEAN DECREASE VARIABLE IMPORTANCE



# OUT-OF-BAG SAMPLES (OOB)

One can show that, on average, drawing  $n$  samples from  $n$  observations with replacement results in about  $2/3$  of the observations being selected.

The remaining one-third of the observations not used are referred to as **out-of-bag (OOB)**

# OUT-OF-BAG SAMPLES (OOB)

We can think of it as a for-free **cross-validation**

The observations that aren't included serve as **test** data

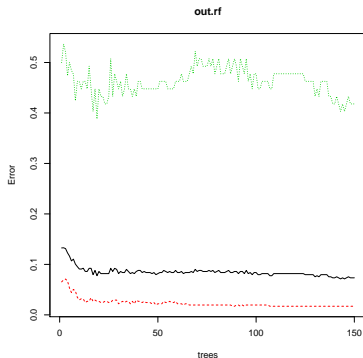
This provides a free estimate of prediction risk for each tree

We can therefore get an overall estimate of prediction risk by averaging these estimates over **all** bootstrapped trees

(Same idea applies here as for getting a prediction at a test  $X$ . For each  $X_i$ , we can take a majority vote of all the classifications when  $(X_i, Y_i)$  is OOB)

# OUT-OF-BAG SAMPLES (OOB)

We can use the OOB samples to choose the number of trees  $B$  to consider



As we are taking an average, we can iteratively compute small batches, stopping when OOB error rate stabilizes

# PERMUTATION VARIABLE IMPORTANCE

Consider the  $b^{th}$  bootstrap sample

1. The OOB prediction accuracy is recorded
2. Then, the  $j^{th}$  feature is randomly permuted in the OOB samples
3. The prediction error is recomputed and the change in prediction error is recorded

**INTUITION:** If a feature is highly important, then the OOB prediction error should increase substantially after permuting the OOB values for that feature

# PROXIMITY

Choose any two observations on the training data:  $i, i'$

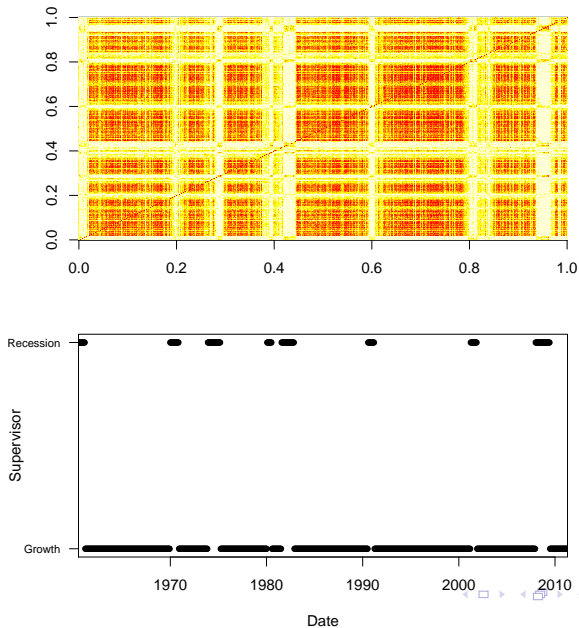
We can record

$$\text{proximity}(i, i') = \frac{\# \text{ times } i, i' \text{ are in the same leaf}}{\# \text{ times } i, i' \text{ occur in same tree}}$$

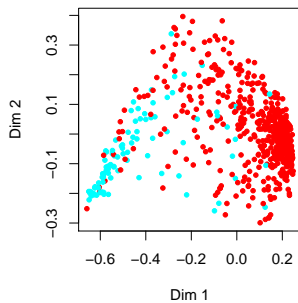
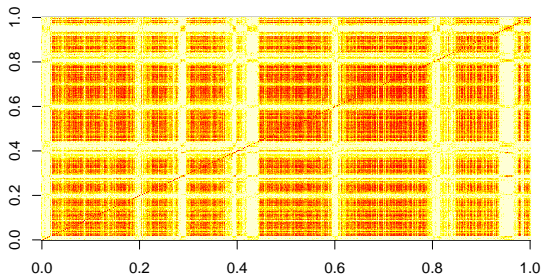
Values near 1 indicate "close" observations and values near 0 indicate "far" observations



# PROXIMITY



# PROXIMITY



# PROXIMITY

```
#First image
out.rf = randomForest(X,Y,proximity=TRUE)
par(mfrow=c(2,1),mar=c(4,4,1.2,4))
image(1 - out.rf$proximity)
plot(dates, Y,
      xlab='Date', ylab='Supervisor')
dev.off()
```

```
#Second image
out.rf = randomForest(X,Y,proximity=TRUE)
par(mfrow=c(2,1),mar=c(4,4,1.2,4))
image(1 - out.rf$proximity)
MDSplot(out.rf,Y,palette=rainbow(2))
```

# Random forest

# RANDOM FOREST

Random Forest is a small extension of Bagging, in which we attempt to **decorrelate** the bootstrap trees

**IDEA:** Draw a bootstrap sample and start to build a tree

- At each split, we randomly select  $m$  of the possible  $p$  features as candidates for the split.
- A new sample of size  $m$  of the features is taken at each split.

Usually  $m = \sqrt{p}$  for classification and  $p/3$  for regression  
(this would be 7 out of 56 features for GDP data)

In other words, at each split, we aren't even allowed to consider the majority of possible features!

# RANDOM FOREST

## What is going on here?

Suppose there is 1 really strong feature and many mediocre ones.

- Then each tree will have this one feature in it,
- Therefore, each tree will look very **similar** (i.e. highly correlated).
- Averaging highly correlated things leads to much less variance reduction than if they were uncorrelated.

If we don't allow some trees/splits to use this important feature, each of the trees will be much less similar and hence much less correlated.

Bagging is Random Forest when  $m = p$

(That is, when we can consider all the features at each split)

# RANDOM FOREST

An average of  $B$  uncorrelated random variables has variance

$$\frac{\sigma^2}{B}$$

An average of  $B$  random variables has variance

$$\rho\sigma^2 + \frac{(1 - \rho)\sigma^2}{B}$$

for correlation  $\rho$

As  $B \rightarrow \infty$ , the second term goes to zero, but the first term remains

Hence, correlation of the trees limits the benefit of averaging

# RANDOM FOREST

Another way to decorrelate the trees is by introducing **noise features**

Generate a few new features (say  $0.01p$ ) that are not related to the supervisor

In some bootstrap samples, this feature will be included in the tree, adding a decorrelating effect



# RANDOM FOREST: BIAS AND VARIANCE

With either approach, we are trading **bias** and **variance** again

Bagging has the same bias as the underlying procedure, but may not get much variance reduction

Random Forest is biased due to subsampling/noise features, but gets more variance reduction by decreasing  $\rho$

(recall that the variance is  $\rho\sigma^2 + \frac{(1-\rho)\sigma^2}{B}$ )

# Evaluating Classifications

# SENSITIVITY AND SPECIFICITY

**SENSITIVITY:** The proportion of times we label **recession**, given that **recession** is the correct answer.

**SPECIFICITY:** The proportion of times we label **no recession**, given that **no recession** is the correct answer.

We can think of this in terms of hypothesis testing. If

$H_0$  : no recession,

then

**SENSITIVITY:**  $P(\text{reject } H_0 | H_0 \text{ is false}), [1 - P(\text{Type II error})]$

**SPECIFICITY:**  $P(\text{accept } H_0 | H_0 \text{ is true}), [1 - P(\text{Type I error})]$

# CONFUSION MATRIX

We can report our results in a matrix:

		Truth	
		Recession	No Recession
Our Predictions	Recession	(A)	(B)
	No Recession	(C)	(D)

The total number of each combination is recorded in the table.

The overall miss-classification rate is

$$\frac{(B) + (C)}{(A) + (B) + (C) + (D)} = \frac{(B) + (C)}{\text{total observations}}$$

What is the sensitivity/specificity?

# CONFUSION MATRIX

We can report our results in a matrix:

		Truth	
		Recession	No Recession
Our Predictions	Recession	(A)	(B)
	No Recession	(C)	(D)

The total number of each combination is recorded in the table.

The overall miss-classification rate is

$$\frac{(B) + (C)}{(A) + (B) + (C) + (D)} = \frac{(B) + (C)}{\text{total observations}}$$

What is the sensitivity/specificity?

(Sensitivity is  $(A)/[(A) + (C)]$ , Specificity is  $(D)/[(B) + (D)]$ )

# TREE RESULTS: CONFUSION MATRICES

			Truth		Mis-Class
			Growth	Recession	
Our Preds	NULL	Growth	111	26	18.9%
		Recession	0	0	
	TREE	Growth	99	3	10.9%
		Recession	12	23	
	RANDOM FOREST	Growth	102	5	10.2%
		Recession	9	21	
	BAGGING	Growth	104	3	7.3%
		Recession	7	23	

# TREE RESULTS: SENSITIVITY & SPECIFICITY

	Sensitivity	Specificity
NULL	0.000	1.000
TREE	0.884	0.891
RANDOM FOREST	0.807	0.918
BAGGING	0.884	0.936

# OUT-OF-BAG ERROR ESTIMATION FOR BAGGING

		Truth		Miss-Class
		Growth	Recession	
OOB BAGGING	Growth	401	9	6.71%
	Recession	23	44	
TEST BAGGING	Growth	104	3	7.3%
	Recession	7	23	



# RANDOM FOREST IN R

```
require(randomForest)
out.rf = randomForest(X,Y,importance=TRUE,mtry=ncol(X))
class.rf = predict(out.rf,X_0)
```

## NOTES:

- The **importance** statement tells it to produce the variable importance measures
- the **mtry = ncol(X)** tells **randomForest** to consider all the features at each split

(This particular choice corresponds to bagging. Leaving this out uses the default  $\sqrt{p}$ )

- **randomForest** also supports formulae

```
out.rf = randomForest(Y~.,data=X)
```

However, it can take much longer to run

# RANDOM FOREST IN R

Call:

```
randomForest(x = X, y = Y, mtry = ncol(X), importance=T)
```

```
      Type of random forest: classification
```

```
      Number of trees: 500
```

```
      No. of variables tried at each split: 56
```

```
      OOB estimate of  error rate: 7.17%
```

Confusion matrix:

```
      0  1 class.error
```

```
0 401  9  0.02195122
```

```
1  25 42  0.37313433
```

# RANDOM FOREST IN R

```
> head(importance(out.rf,type=1))#Permutation
      MeanDecreaseAccuracy
Alabama                3.7277511
Alaska                 1.7941463
Arizona                2.9659623
Arkansas               -0.8341577
California              7.2973572
> head(importance(out.rf,type=2))#Mean decrease
      MeanDecreaseGini
Alabama                0.4551073
Alaska                 1.6440170
Arizona                0.7025527
Arkansas               0.3503138
California             1.4616203
#variable importance plot:
varImpPlot(out.rf,type=2)
```

# MISSING DATA/IMPUTATION

In practice, there will often be **missing data**

Estimating this missing data is known as **imputation**

**RANDOM FOREST** provides a method for imputation

It follows two steps:

1. **na.roughfix**: uses either the median or mode to impute missing values  
$$X \leftarrow \text{na.roughfix}(X)$$
2. **rf.impute**: Gets the proximity matrix, and re-computes the imputation
  - ▶ For numeric features, it uses weighted (with respect to proximity) average
  - ▶ For categorical features, it uses the category that maximizes the average proximity

(Both of these only use the original, non-missing observations)

# MISSING DATA/IMPUTATION

```
require(randomForest)
x = rnorm(12)
xNA = x
xNA[sample(12,2,replace=F)] = NA
X = matrix(xNA,nrow=6)
Y = matrix(x,nrow=6) %*% c(1,.5) + rnorm(6)
Xnew = rfImpute(X,Y)
```

# MISSING DATA/IMPUTATION

```
> X
```

```
      NA -1.2589350282
0.6202015 0.1780122216
-0.9340213 -0.6483047015
0.1142546 0.0489260332
-1.1039581      NA
0.2064204 0.0007548357
> matrix(x,nrow=6)
0.1485857 -1.2589350282
0.6202015 0.1780122216
-0.9340213 -0.6483047015
0.1142546 0.0489260332
-1.1039581 -0.2468085986
0.2064204 0.0007548357
```

```
> Xnew
```

```
0.1287655 -1.2589350282
0.6202015 0.1780122216
-0.9340213 -0.6483047015
0.1142546 0.0489260332
-1.1039581 -0.2015763074
0.2064204 0.0007548357
```

# Additional random forest topics

# BAGGING THEORY

Suppose we are using random forest for regression

The bootstrap estimator is:

$$\hat{f}_{\text{bag}}(X) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b^*(X)$$

For bootstrap trees:

$$\hat{f}_1^*(X), \dots, \hat{f}_B^*(X)$$

To see what is really going on, we need some bootstrap theory



# BOOTSTRAP THEORY

Suppose we have i.i.d. observations  $Z_i$  from  $\mathbb{P}$

(Call the data  $\mathcal{D}$ )

Suppose we want to know the distribution of some procedure that is a function of  $\mathcal{D}$ ,  $\hat{\theta}(\mathcal{D})$

$$\mathbb{P}(\hat{\theta}(\mathcal{D}) \in A)$$

This is unknown and can be estimated with a **plug-in** estimator of the measure  $\mathbb{P}$

- **PARAMETRIC:**

$$\mathbb{P}_{\hat{\theta}}(\hat{\theta}(\mathcal{D}) \in A)$$

- **NONPARAMETRIC:**

$$\hat{\mathbb{P}}(\hat{\theta}(\mathcal{D}) \in A)$$

# BOOTSTRAP THEORY

In either case, **computing** with respect to  $\mathbb{P}_{\hat{\theta}}$  or  $\hat{\mathbb{P}}$  is commonly infeasible but **sampling** from them is easy

**EXAMPLE:** Estimate the statistical functional  $\theta(\mathbb{P}) = \int z d\mathbb{P}(z) = \mathbb{E}Z$  with the plug-in estimator  $\theta(\hat{\mathbb{P}}) = \bar{Z}$ .

Compute the mean of  $\theta(\hat{\mathbb{P}})$  w.r.t  $\hat{\mathbb{P}}$ :

$$\hat{\mathbb{P}}\theta(\hat{\mathbb{P}})$$

Let  $\mathcal{D} = \{4, 10\}$ . Then there are 4 possibilities (in general  $n^n$ )

$$\hat{\mathbb{P}}\theta(\hat{\mathbb{P}}) = \frac{1}{4} \left( \frac{1}{2}(4 + 4) + \frac{1}{2}(10 + 4) + \frac{1}{2}(4 + 10) + \frac{1}{2}(10 + 10) \right) = 7$$

# BOOTSTRAP THEORY

Due to the extremely fast growth of  $n^n$  or intractability of  $\mathbb{P}_{\hat{\theta}}$ , we need to approximate the distribution

→ replace them with a Monte-Carlo approximation

Draw  $\mathcal{D}^* = Z_1^*, Z_2^*, \dots, Z_B^*$  from  $\hat{\mathbb{P}}$  or  $\mathbb{P}_{\hat{\theta}}$ :

$$\hat{\mathbb{P}}(\hat{\theta}(\mathcal{D}) \in A) \text{ replace with } \frac{1}{B} \sum_{b=1}^B \mathbf{1}(\hat{\theta}(\mathcal{D}^*) \in A)$$

**IMPORTANT QUESTION:** How good is this approximation?

# BOOTSTRAP THEORY

Write  $\mathbb{P}^*$  as the bootstrap distribution.

Suppose we want to show some result about

$$\sqrt{n}(\theta(\mathbb{P}^*) - \theta(\mathbb{P})) = \sqrt{n}(\theta(\mathbb{P}^*) - \theta(\hat{\mathbb{P}})) + \sqrt{n}(\theta(\hat{\mathbb{P}}) - \theta(\mathbb{P}))$$

**EXAMPLE:** If  $\mathbb{P}Z^2 < \infty$ , then

$$\sqrt{n}(\theta(\hat{\mathbb{P}}) - \theta(\mathbb{P})) = \sqrt{n}(\bar{Z} - \mu) \xrightarrow{D} N(0, \sigma^2)$$

It turns out that

$$\sqrt{n}(\bar{Z}^* - \bar{Z}) \xrightarrow{D} N(0, \sigma^2)$$

as well

# ADDITIONAL RANDOM FOREST TOPICS

**CLAIM:** Random forest cannot overfit.

This is and isn't true. Write

$$\hat{f}_{rf}^B(X) = \frac{1}{B} \sum_{b=1}^B T(X; \Theta_b)$$

( $\Theta_b$  are all the feature split points and terminal node values of  $b^{th}$  tree)

Increasing  $B$  does not cause Random forest to overfit, rather removes the Monte-Carlo-like approximation error

$$\hat{f}_{rf}(X) = \mathbb{E}_{\Theta} T(X, \Theta) = \lim_{B \rightarrow \infty} \hat{f}_{rf}^B(X)$$

However, **this limit can overfit the data**, the average of fully grown trees can result in too complex of a model

(Note that Segal (2004) shows that a small benefit can be derived by stopping each tree short, but thus induces another tuning parameter)

# ADDITIONAL RANDOM FOREST TOPICS

- Leo Breiman advocated for a “linear combination” random forest. At each node, random linear combinations of subselected features with weights in  $[-1, 1]$  are split upon. He found this approach improved upon the standard random forest.
- Introduction of noise features improves performance
- Using **subsampling** instead of bootstrap to generate trees
- Adaptive nearest neighbors via proximity scores