


1 What is a (decision) tree?



Basically...

- Trees are models that *stratify* or *segment* the predictor space into a number of simple regions.
- We predict all observations in a region the same prediction
- The three regions R1, R2, and R3 are the *terminal nodes*

Main takeaways...

- Trees are simple and useful for interpretation.
- Basic trees are not great at prediction.
- More modern methods that use trees are much better.

2 How do we build a tree?

1. Divide the predictor space into M non-overlapping regions R_1, \dots, R_M
-this is done via greedy, recursive, binary splitting
2. Every observation that falls into a given region R_m is given the same prediction
 - **Regression:** The average of the responses for a region
 - **Classification:** Determined by majority (or plurality) vote in that region

2.1 Important:

- Trees can only make rectangular regions that are *aligned* with the coordinate axis.
- The fit is *greedy*, which means that after a split is made, all further decisions are conditional on that split.

3 Regression Trees

For a given partition R_1, \dots, R_M , the model for the response is

$$f(x) = \sum_{m=1}^M c_m \mathbf{1}_{R_m}(x)$$

For squared error loss, if $n_m = \sum_{i=1}^n \mathbf{1}_{R_m}(X_i)$, then

$$\hat{c}_m = n_m^{-1} \sum_{i: x_i \in R_m} Y_i$$

As for the regions, M encodes the tree complexity

This is challenging as considering all possible regions is computationally *infeasible*

(This would involve sifting through all $M \leq n$ and all configurations for R_m .)

3.1 Model Selection for Trees

As a greedy approximation, do the following

1. **Grow a large tree:** T_{\max} , stopping when some minimal terminal node size requirement is met
2. **Cost-complexity pruning:** For all $\lambda \geq 0$

$$C_\lambda(T) = \sum_{m=1}^M \sum_{i: x_i \in R_m} (Y_i - \hat{c}_m)^2 + \lambda M$$

(Note that often it is written that $|T| = M$)

3. **Weakest link pruning:** For each λ , there is a unique smallest T_λ that minimizes $C_\lambda(T)$. Eliminating nodes that produce the smallest increase in training error produces a sequence of solutions that must contain T_λ
(many details omitted)

4 Classification Trees

The only modification for *classification* is choice of *loss function*

For region m and class g , we get training proportions

$$\hat{p}_{mg}(x) = \mathbf{1}_{R_m}(x) n_m^{-1} \sum_{i: X_i \in R_m} \mathbf{1}(Y_i = g)$$

Our classification is

$$\hat{g}(x) = \max_g \hat{p}_{mg}(x)$$

4.1 And measuring quality of fit?

Different measures of *node impurity* (loss function in tree terminology)

| | |
|-----------------------------------|--|
| classification error rate: | $E = 1 - \max_k(\hat{p}_{mk})$ |
| Gini index: | $G = \sum_k \hat{p}_{mk}(1 - \hat{p}_{mk})$ |
| cross-entropy: | $D = - \sum_k \hat{p}_{mk} \log(\hat{p}_{mk})$ |

Both G and D can be thought of as measuring the purity of the classifier (small if all \hat{p}_{mk} are near zero or 1). These are preferred over the classification error rate.

(Also, E isn't differentiable and hence not as amenable to numerical optimization)

We build a classifier by *growing* a tree that minimizes G or D .

5 Advantages and Disadvantages of Trees:

- + Trees are very easy to explain (much easier than even linear regression).
- + Some people believe that decision trees mirror human decision.
- + Trees can easily be displayed graphically no matter the dimension of the data.
- + Trees can easily handle qualitative predictors without the need to create dummy variables.
- Trees aren't very good at prediction.

To fix this last one, we can try to grow many trees and average their performance. This can be done through *Bagging!* *Boosting* -and- *Random Forests* are variations which also address this last issue.

Trees, in any of their varied forms of aggregation as well, can overfit and thus can be “pruned.” Below is an illustration of this procedure:

